Metrics see the person or reject averaging; the shrunk true person (green) passes almost all of them

A. Does the metric see the person and reject subtle averaging?

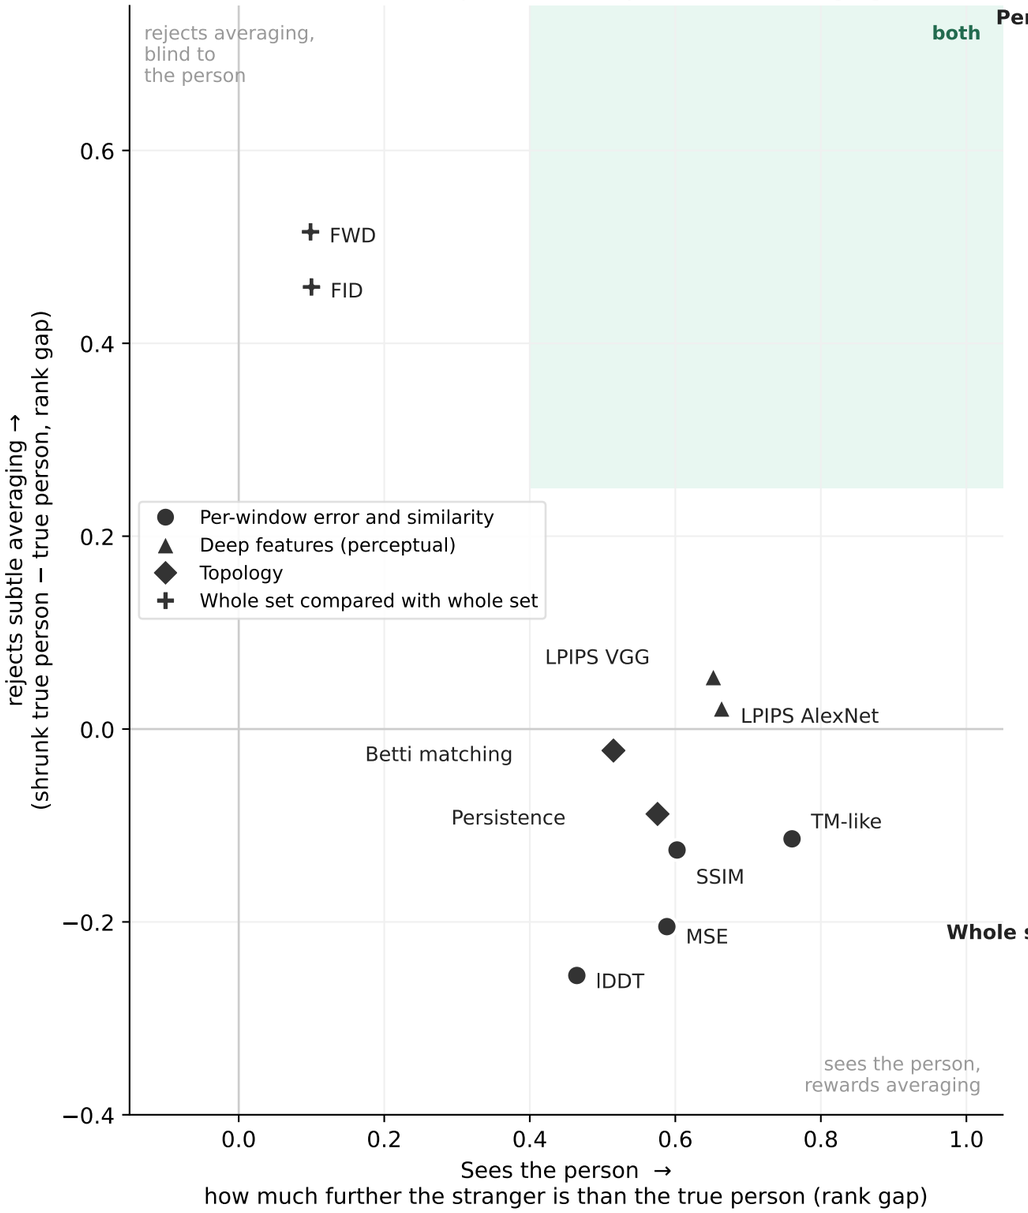

B. How far each metric puts each fill, in units of a stranger (0 = the target itself, 1 = a stranger)

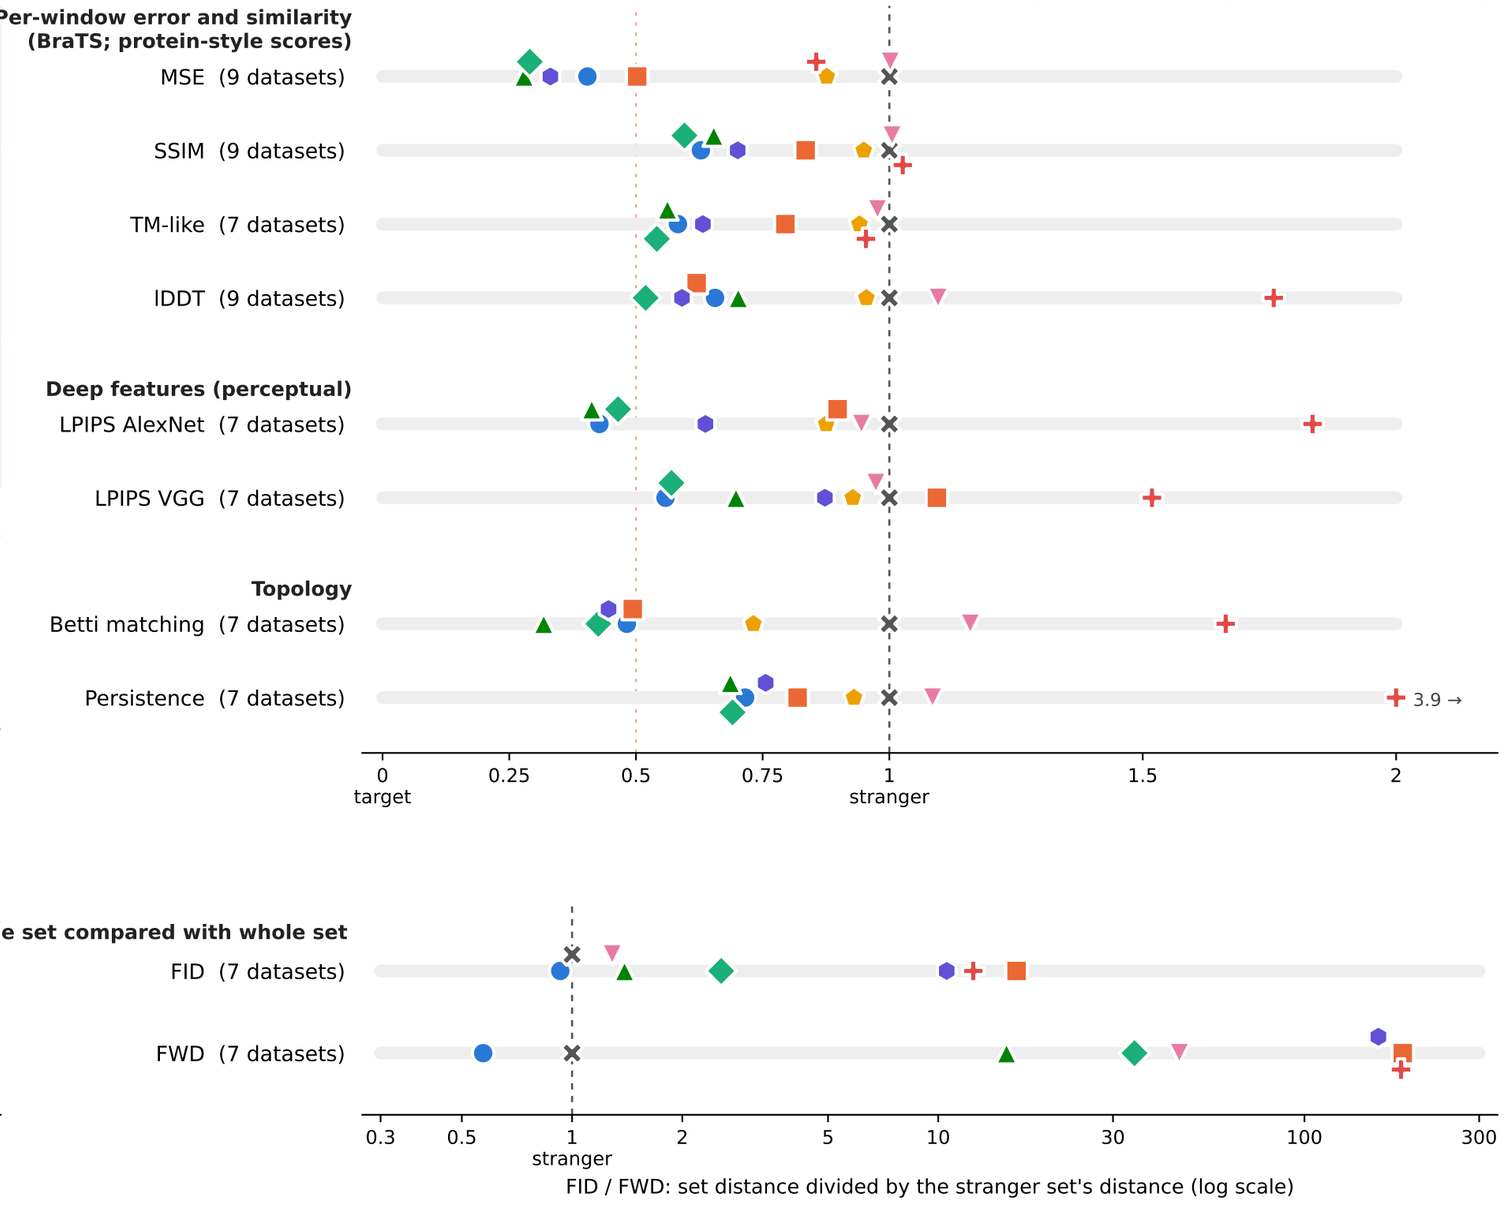

Ladder:
true person (test-retest) · oracle look-alike* · sibling (QTIM, IMDb-WIKI) · dataset mean · stranger

Checks:
true person shrunk to the mean · mean of k look-alikes* · wrong detail · white noise

A: per metric, from its ordering of the ladder (rank gaps / ladder length, mean over datasets): stranger minus true person (x), true person shrunk to the mean minus true person (y); shaded: x >= 0.4 and y >= 0.25.
B: R = (d(fill) - d(target)) / (d(16 strangers) - d(target)), mean over datasets; per-image metrics on a linear axis, FID/FWD on a log axis.
* Look-alikes: oracle = chosen with the hidden truth (upper bound); mean of k look-alikes (k = 16), all nine sets. True person = repeat scan, or the other side where there is no repeat (PAPILA, IXI).
Brain LPIPS, Betti and persistence (and IXI Betti) are on other anchors and left out of B's means. Brain is excluded from FID/FWD (its repeat set is about 10x smaller than its stranger set; FID is biased upward for small sets), and set-level rungs with under half the stranger set's windows are left out.
Data: QTIM and IXI brain, CelebA, IMDb-WIKI, GRAPE, PAPILA, UTIRIS NIR/VIS, UPOL; brain, fundus, iris pre-QC. Source: results/cross_modality/common_all.json, values_generative.json, ranks_all.json, ranks_generative.json.