Metrics see the person or reject averaging; the shrunk true person (green) passes almost all of them

A. Crude averaging: is the dataset mean pushed away?

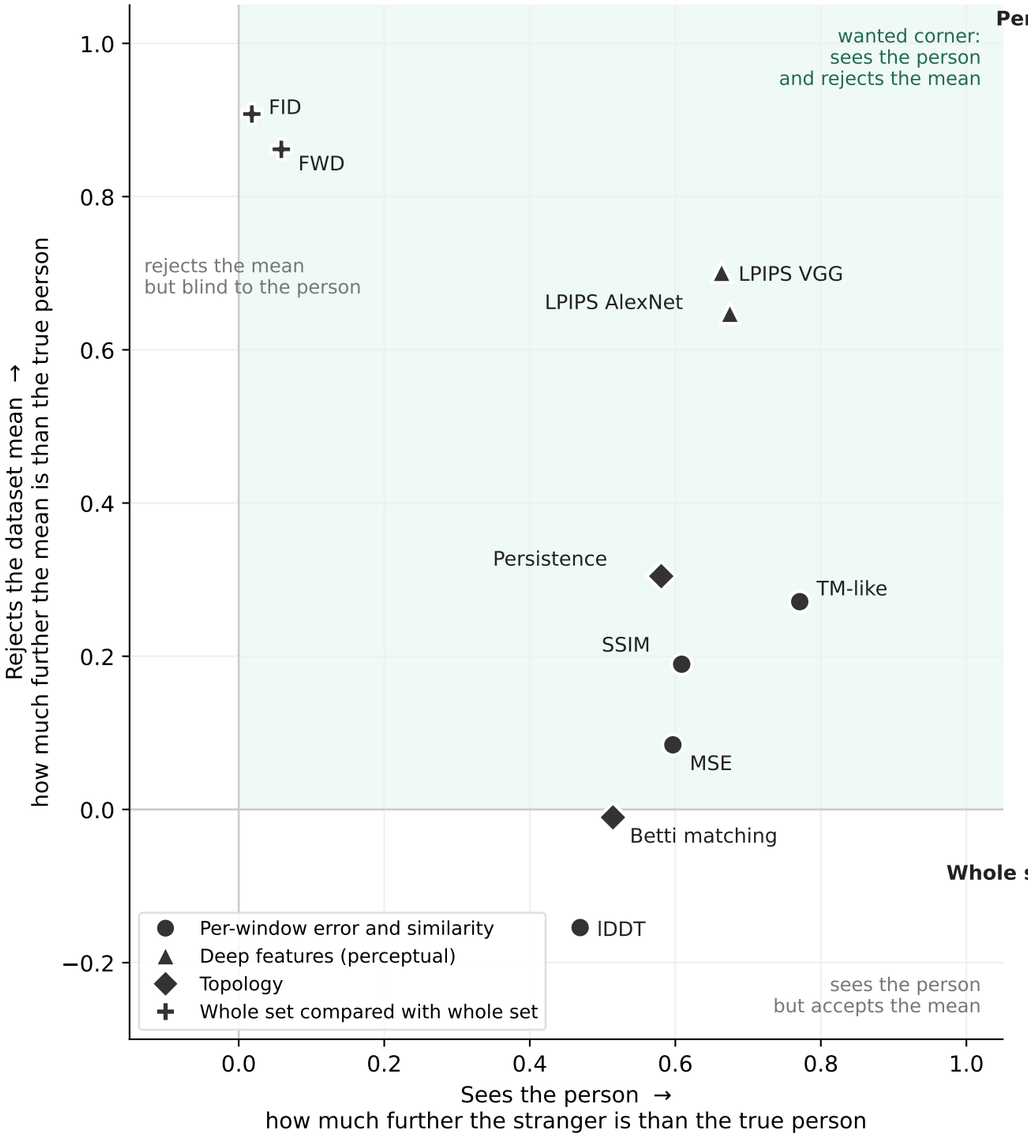

B. How far each metric puts each fill, in units of a stranger (0 = the target itself, 1 = a stranger)

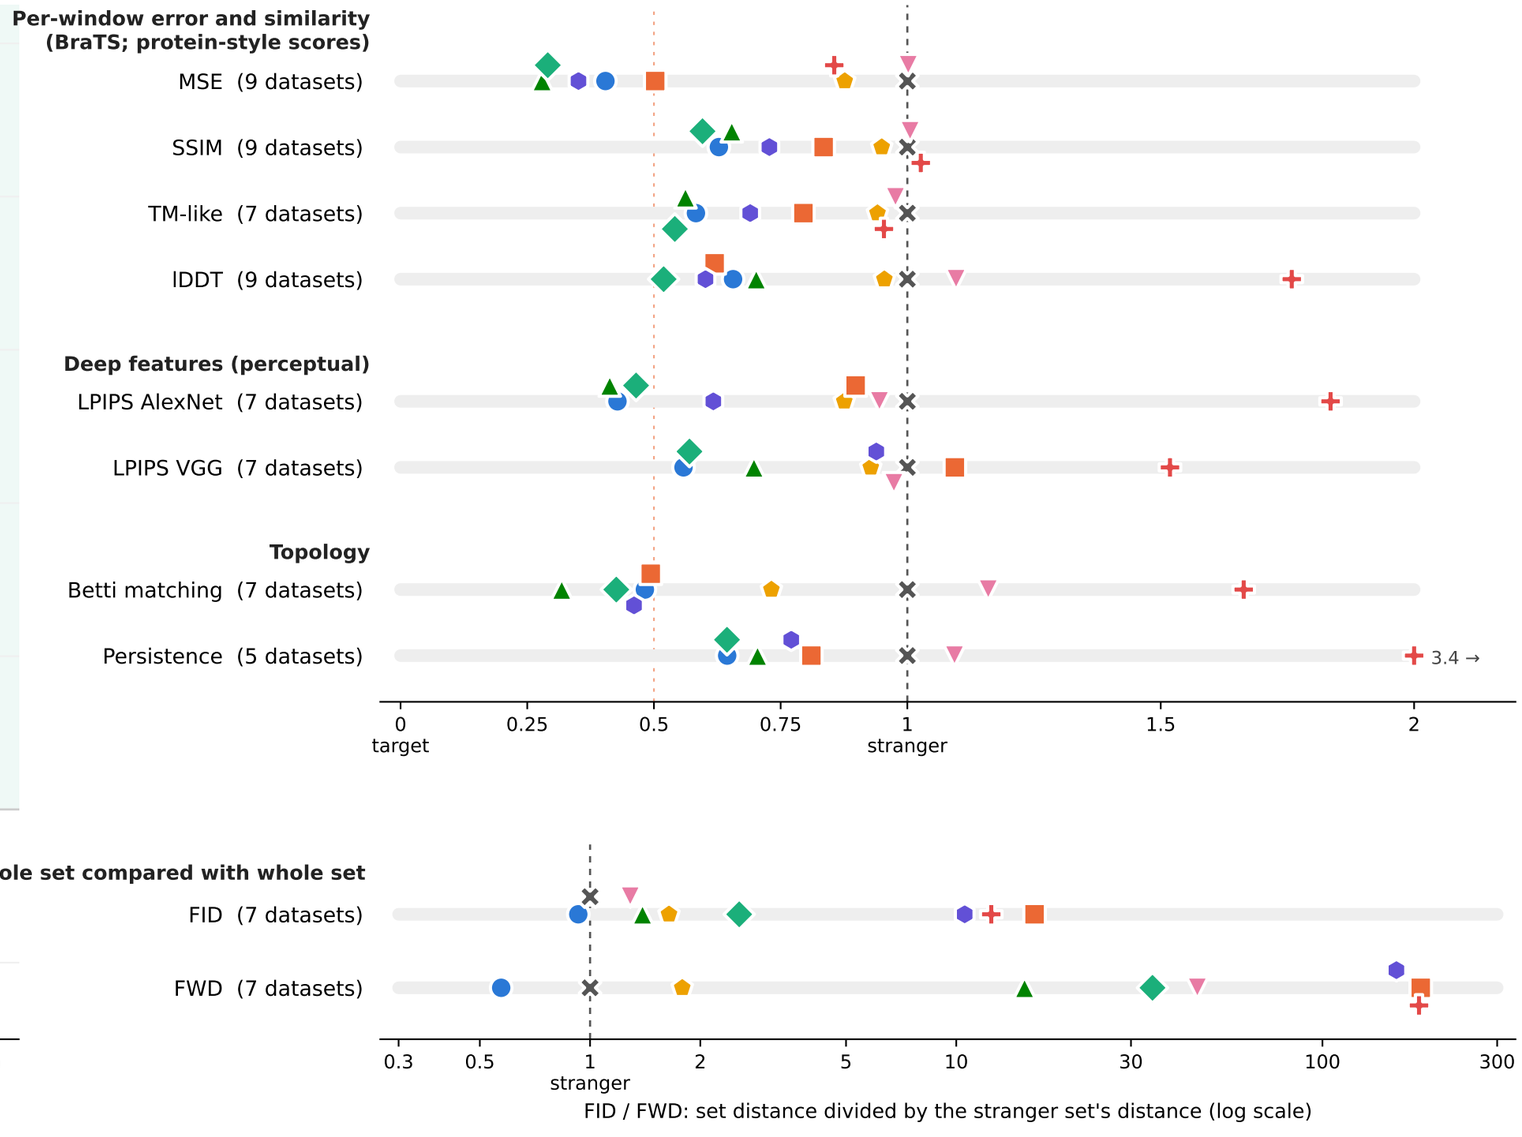

Ladder:
true person (test-retest) · oracle look-alike* · sibling (QTIM, IMDb-WIKI) · dataset mean · stranger

Checks:
true person shrunk to the mean · mean of k look-alikes* (not faces) · wrong detail · white noise

A: per metric, from its ordering of the ladder (rank gaps / ladder length, mean over datasets): stranger minus true person (x), dataset mean minus true person (y). B: R = (d(fill) - d(target)) / (d(16 strangers) - d(target)), mean over datasets; per-image metrics on a linear axis, FID/FWD on a log axis.
* Look-alikes: oracle = chosen with the hidden truth (upper bound); mean of k look-alikes (k = 16) on brain, fundus and iris (not faces). True person = repeat scan, or the other side where there is no repeat (PAPILA, IXI). Brain LPIPS and persistence are repeat-anchored and left out of B's means.
Data: QTIM and IXI brain, CelebA, IMDb-WIKI, GRAPE, PAPILA, UTIRIS NIR/VIS, UPOL; brain, fundus, iris pre-QC. Source: results/cross_modality/common_all.json, values_generative.json, ranks_all.json, ranks_generative.json.